

Project Title: Lamma – Development

Objective

To design and develop an open-source chatbot using the **Llama 3** model family, capable of answering both general and domain-specific queries through Retrieval-Augmented Generation (RAG).

The system will run **entirely on local machines** with a Python-based user interface.

Tasks

1. Model & Environment Setup

- Study the **Llama 3** architecture, model variants, and license terms.
- Install and configure **Ollama**, **vLLM**, or **llama.cpp** for local model serving.
- Run sample prompts to verify local inference.
- Document installation steps, dependencies, and environment configuration.

2. API Layer Development

- Build a **FastAPI** backend to handle user queries and forward them to the model.
- Create REST endpoints such as `/chat` and `/v1/chat/completions`.
- Implement basic authentication (API key or local token).
- Prepare small Python test scripts or Postman collections to validate API calls.

3. Retrieval-Augmented Generation (RAG) Integration

- Set up a **local vector database** (pgvector / Qdrant / Chroma).
- Develop a document ingestion pipeline (PDF → text → chunks → embeddings).
- Integrate retrieval into the chatbot flow (retrieve → context → generate).

- Test the chatbot's ability to answer from uploaded documents.
- Measure accuracy and latency in local setup.

4. Python-Based Front-End Interface

- Build an interactive chatbot interface
- Connect the interface to the FastAPI backend for real-time chat.
- Add features such as "Clear Chat," "Upload Document," "Show References," and "Toggle System Prompt."
- Ensure smooth interaction and proper error handling within the Python interface.

5. Monitoring & Logging

- Implement **Python logging** to record queries, responses, and system errors.
- Track response times and token usage locally.
- Maintain version logs for model, embeddings, and configuration changes.
- Visualize simple metrics (e.g., requests per session, response delay) using Matplotlib or Plotly if needed.

6. Safety & Governance Layer

- Integrate **Llama Guard** or a rule-based filter for input/output moderation.
- Block unsafe, irrelevant, or sensitive queries.
- Display a disclaimer and usage policy within the Python interface.
- Document moderation test cases and outcomes.

7. Documentation & Reporting

- Prepare detailed technical documentation covering:
 - Architecture diagram
 - API structure

- RAG workflow
 - Front-end interface overview
- Include setup and execution instructions for local machines.
 - Prepare a short demo video showing the chatbot in action.
 - Publish source code and documentation in a local or GitHub repository.
-

Deliverables

- Locally running chatbot accessible via Python interface
- Functional API connected to Llama 3 model
- Document-aware (RAG-enabled) query handling
- Logging and moderation mechanisms
- Full technical documentation and demo video